

```
!pip install pyspellchecker
```

```
Collecting pyspellchecker
  Downloading pyspellchecker-0.8.2-py3-none-any.whl.metadata (9.4 kB)
  Downloading pyspellchecker-0.8.2-py3-none-any.whl (7.1 MB)
    7.1/7.1 MB 119.3 MB/s eta 0:00:00
Installing collected packages: pyspellchecker
Successfully installed pyspellchecker-0.8.2
```

```
!pip install contractions
```

```
Collecting contractions
  Downloading contractions-0.1.73-py2.py3-none-any.whl.metadata (1.2 kB)
Collecting textsearch>=0.0.21 (from contractions)
  Downloading textsearch-0.0.24-py2.py3-none-any.whl.metadata (1.2 kB)
Collecting anyascii (from textsearch>=0.0.21->contractions)
  Downloading anyascii-0.3.2-py3-none-any.whl.metadata (1.5 kB)
Collecting pyahocorasick (from textsearch>=0.0.21->contractions)
  Downloading pyahocorasick-2.1.0-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (13 kB)
  Downloading contractions-0.1.73-py2.py3-none-any.whl (8.7 kB)
  Downloading textsearch-0.0.24-py2.py3-none-any.whl (7.6 kB)
  Downloading anyascii-0.3.2-py3-none-any.whl (289 kB)
    289.9/289.9 kB 22.4 MB/s eta 0:00:00
  Downloading pyahocorasick-2.1.0-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (118 kB)
    118.3/118.3 kB 8.5 MB/s eta 0:00:00
Installing collected packages: pyahocorasick, anyascii, textsearch, contractions
Successfully installed anyascii-0.3.2 contractions-0.1.73 pyahocorasick-2.1.0 textsearch-0.0.24
```

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import re
import nltk
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
from nltk.stem import WordNetLemmatizer
from contractions import fix
```

```
df = pd.read_csv('/content/Train.csv')
df.head()
```

```

text  label
0    I grew up (b. 1965) watching and loving the Th...    0
1    When I put this movie in my DVD player, and sa...    0
2    Why do people who do not know what a particula...    0
3    Even though I have great interest in Biblical ...    0
4    Im a die hard Dads Army fan and nothing will e...    1
```

```
df.info()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 40000 entries, 0 to 39999
Data columns (total 2 columns):
 #   Column  Non-Null Count  Dtype
---  -
 0   text    40000 non-null    object
 1   label   40000 non-null    int64
dtypes: int64(1), object(1)
memory usage: 625.1+ KB
```

```
df['text'][1]
```


```
'When I put this movie in my DVD player, and sat down with a coke and some chips, I had some expectations. I was hoping that this movie would contain some of the strong-points of the first movie: Awsome animation, good flowing story, excellent voice cast, funny comedy and a kick-ass soundtrack. But, to my disappointment, not any of this is to be found in Atlantis: Milo's Return. Had I read some reviews first, I might not have been so let down. The following paragraph will be directed to those who have seen the first movie, and who enjoyed it primarily for the points mentioned.<br /><br />When the first scene appears, your in for a shock if you just picked Atlantis: Milo's Return from the displav-case at your local videoshon (or whatever). and had the expektations I had. The music feels as a had imita
```

```
plt.figure(figsize=(10,4))
df['review_length_char'] = df['text'].apply(len)
```

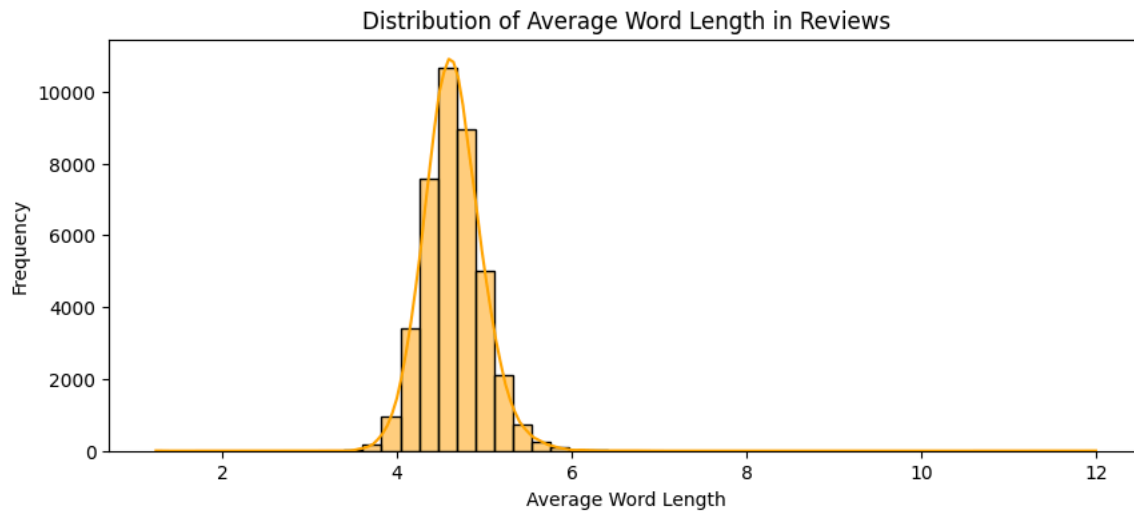
```
sns.histplot(df['review_length_char'], bins=50, kde=True, color='blue')
plt.title("Distribution of Review Length in Characters")
plt.xlabel("Number of Characters")
plt.ylabel("Frequency")
plt.show()
```



```
plt.figure(figsize=(10,4))
df['review_length_word'] = df['text'].apply(lambda x: len(x.split()))
sns.histplot(df['review_length_word'], bins=50, kde=True, color='green')
plt.title("Distribution of Review Length in Words")
plt.xlabel("Number of Words")
plt.ylabel("Frequency")
plt.show()
```



```
df['avg_word_length'] = df['text'].apply(lambda x: np.mean([len(word) for word in x.split()]))
plt.figure(figsize=(10,4))
sns.histplot(df['avg_word_length'], bins=50, kde=True, color='orange')
plt.title("Distribution of Average Word Length in Reviews")
plt.xlabel("Average Word Length")
plt.ylabel("Frequency")
plt.show()
```



```
print(df['label'].unique())
positive_reviews = df[df['label'] == 'pos']['text']
negative_reviews = df[df['label'] == 'neg']['text']
```

[0 1]

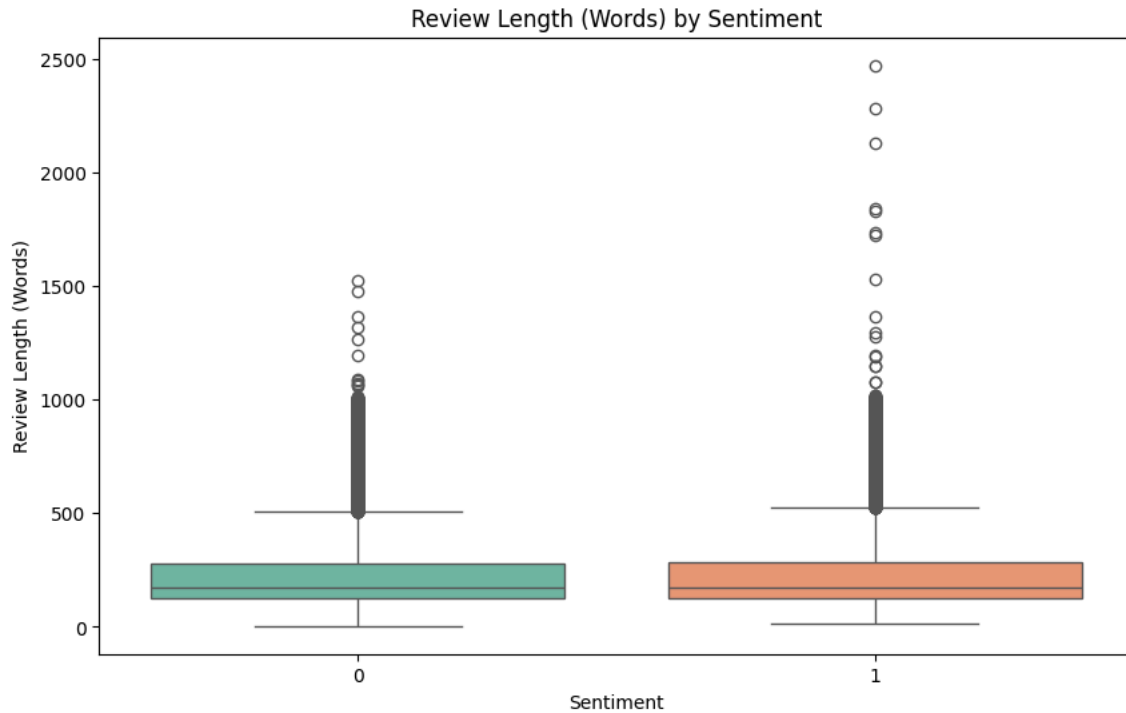
```
plt.figure(figsize=(10,6))

sns.boxplot(x='label', y='review_length_word', data=df, palette="Set2")
plt.title("Review Length (Words) by Sentiment")
plt.xlabel("Sentiment")
plt.ylabel("Review Length (Words)")
plt.show()

plt.figure(figsize=(10,6))
# Changed 'sentiment' to 'label' to use the existing column for sentiment
sns.boxplot(x='label', y='review_length_char', data=df, palette="Set3")
plt.title("Review Length (Characters) by Sentiment")
plt.xlabel("Sentiment")
plt.ylabel("Review Length (Characters)")
plt.show()
```

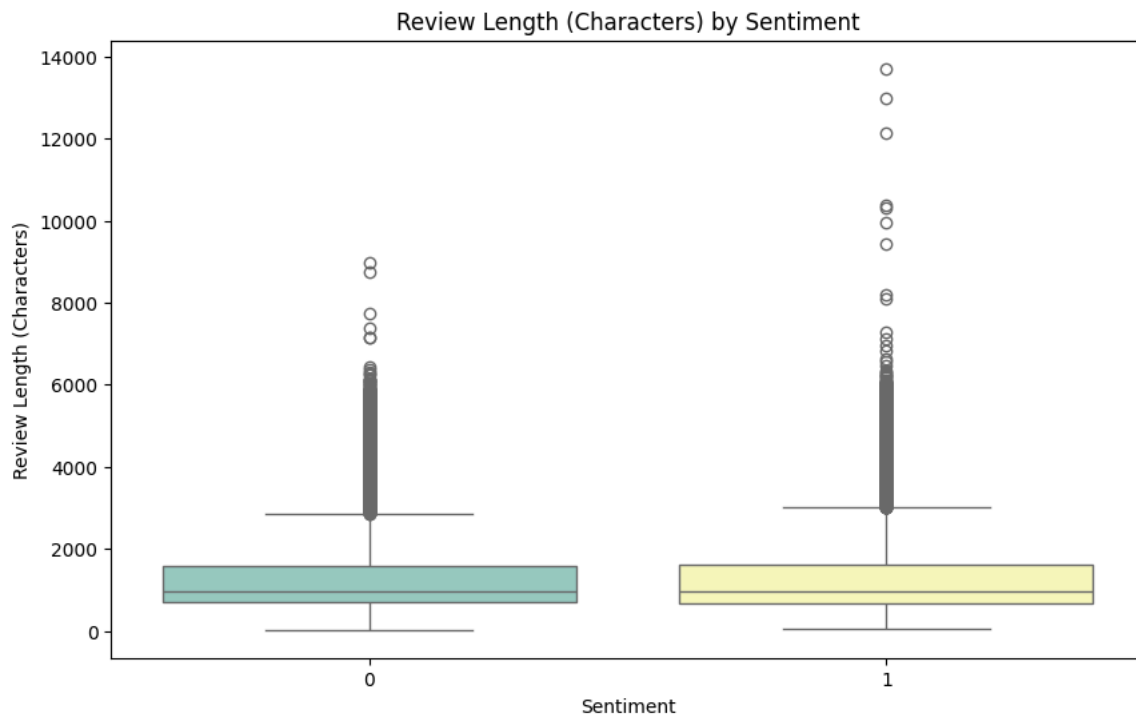
```
<ipython-input-13-5c4a4632a18a>:3: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `leg`
 sns.boxplot(x='label', y='review_length_word', data=df, palette="Set2")



```
<ipython-input-13-5c4a4632a18a>:11: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `leg`
 sns.boxplot(x='label', y='review_length_char', data=df, palette="Set3")



```
import nltk
nltk.download('stopwords')
nltk.download('punkt')
nltk.download('wordnet')
from nltk.stem import WordNetLemmatizer
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
import re
```

```

from spellchecker import SpellChecker
from contractions import fix

lemmatizer = WordNetLemmatizer()
stop_words = set(stopwords.words('english'))

def preprocess_text(text):
    # Lowercase the text
    text = text.lower()
    # Expand contractions
    text = fix(text)

    text = re.sub(r'<.*?>', '', text)
    text = re.sub(r'https?://\S+|www\.\S+', '', text)
    # Remove special characters (keeping words, spaces, hyphens, apostrophes)
    text = re.sub(r"[^\w\s'-]", ' ', text)
    # Tokenize the text
    tokens = word_tokenize(text)
    # Remove stopwords and lemmatize
    tokens = [lemmatizer.lemmatize(word) for word in tokens if word not in stop_words]
    # Join tokens back into a single string
    return ' '.join(tokens)

df['text'] = df['text'].apply(preprocess_text)

[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Unzipping corpora/stopwords.zip.
[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data] Unzipping tokenizers/punkt.zip.
[nltk_data] Downloading package wordnet to /root/nltk_data...

-----
LookupError                                Traceback (most recent call last)
<ipython-input-14-35c036a5aa6d> in <cell line: 0>()
    30     # Join tokens back into a single string
    31     return ' '.join(tokens)
--> 32 df['text'] = df['text'].apply(preprocess_text)

-----
lib.pyx in pandas._libs.lib.map_infer()

/usr/local/lib/python3.11/dist-packages/nltk/data.py in find(resource_name, paths)
    577     sep = "*" * 70
    578     resource_not_found = f"\n{sep}\n{msg}\n{sep}\n"
--> 579     raise LookupError(resource_not_found)
    580
    581

LookupError:
*****
Resource punkt_tab not found.
Please use the NLTK Downloader to obtain the resource:

>>> import nltk
>>> nltk.download('punkt_tab')

For more information see: https://www.nltk.org/data.html

Attempted to load tokenizers/punkt_tab/english/

Searched in:
- '/root/nltk_data'
- '/usr/nltk_data'
- '/usr/share/nltk_data'
- '/usr/lib/nltk_data'
- '/usr/share/nltk_data'
- '/usr/local/share/nltk_data'
- '/usr/lib/nltk_data'
- '/usr/local/lib/nltk_data'
*****

```

df['text'][1]

df['text'].groupby(df['label']).count()

```

↗
text
label
0    20019
1    19981

dtype: int64

```

```
df.groupby('label').describe()
```

```

↗
review_length_char      review_length_word      avg_word_length
count  mean      std      min  25%   50%   75%   max      count  mean      ...  75%   max      count  mean      std
label
0    20019.0  1292.536990  942.220087  32.0  705.0  973.0  1571.0  8969.0  20019.0  229.204606  ...  279.0  1522.0  20019.0  4.623670  0.328
1    19981.0  1328.083279  1032.236721  65.0  690.0  972.0  1621.0  13704.0  19981.0  233.477954  ...  285.0  2470.0  19981.0  4.657509  0.350

```

2 rows × 24 columns

```
df['length'] = df['text'].apply(len)
df.head()
```

```

↗
text  label  review_length_char  review_length_word  avg_word_length  length
0    I grew up (b. 1965) watching and loving the Th...    0          874          151          4.794702    874
1    When I put this movie in my DVD player, and sa...    0          1811          326          4.558282    1811
2    Why do people who do not know what a particula...    0          983          184          4.347826    983
3    Even though I have great interest in Biblical ...    0          351           69          4.101449    351
4    Im a die hard Dads Army fan and nothing will e...    1          983          178          4.528090    983

```

Vectorization

```
from sklearn.feature_extraction.text import CountVectorizer
```

```
bow = CountVectorizer(analyzer='word').fit(df['text'])
```

```
print(len(bow.vocabulary_))
```

```
↗ 92908
```

```
MESS2 = df['text'][1]
print(MESS2)
```

```
↗ When I put this movie in my DVD player, and sat down with a coke and some chips, I had some expectations. I was hoping that this movie w
```

```
MESS2 = bow.transform([MESS2])
print(MESS2)
```

```

↗ <Compressed Sparse Row sparse matrix of dtype 'int64'
   with 167 stored elements and shape (1, 92908)>
  Coords      Values
(0, 2100)      1
(0, 2438)      1
(0, 3495)      1
(0, 3545)      1
(0, 3593)      1
(0, 4083)      7
(0, 4287)      2
(0, 4644)      1
(0, 4846)      1
(0, 5580)      3
(0, 5733)      1
(0, 5928)      1
(0, 6001)      2
(0, 6181)      1
(0, 6578)      1

```

```
(0, 6868)    2
(0, 7938)    3
(0, 8207)    4
(0, 8897)    1
(0, 10945)   4
(0, 11194)   1
(0, 12410)   4
(0, 12510)   1
(0, 13371)   1
(0, 13633)   1
:           :
(0, 82218)   1
(0, 82245)   2
(0, 82262)   20
(0, 82529)   6
(0, 82591)   1
(0, 82642)   2
(0, 82967)   2
(0, 83225)   5
(0, 85169)   1
(0, 85227)   1
(0, 88210)   1
(0, 88378)   1
(0, 88851)   3
(0, 89510)   1
(0, 89717)   1
(0, 89718)   1
(0, 90160)   1
(0, 90214)   2
(0, 90401)   3
(0, 90645)   1
(0, 90964)   2
(0, 91425)   1
(0, 91428)   1
(0, 92197)   6
(0, 92222)   2
```

```
MESS2.shape
```

```
(1, 92908)
```

```
bow.get_feature_names_out()[78]
```

```
'1050'
```

```
mess_bow = bow.transform(df['text'])
```

```
print('shape of sparse matrix : ',mess_bow.shape)
print('ammount of non_zero ocarences: ',mess_bow.nnz)
```

```
shape of sparse matrix : (40000, 92908)
ammount of non_zero ocarences: 5463770
```

```
sparsity = (100.0*mess_bow.nnz/(mess_bow.shape[0]*mess_bow.shape[1]))
print('sparsity : {}'.format(sparsity))
```

```
sparsity : 0.14702097774142162
```

```
from sklearn.feature_extraction.text import TfidfTransformer
```

```
TF = TfidfTransformer().fit(mess_bow)
```

```
TF.idf_[bow.vocabulary_['bad']]
```

```
np.float64(2.4440544114586342)
```

```
df['label'] = df['label'].map({'pos': 1, 'neg': 0})
```

```
df['label']
```



	label
0	NaN
1	NaN
2	NaN
3	NaN
4	NaN
...	...
39995	NaN
39996	NaN
39997	NaN
39998	NaN
39999	NaN

40000 rows × 1 columns

dtype: float64

```

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(df['text'], df['label'])

from sklearn.feature_extraction.text import TfidfVectorizer, CountVectorizer # Import CountVectorizer here as well

count_vectorizer = CountVectorizer(max_features=5000)

tfidf_vectorizer = TfidfVectorizer(max_features=5000)

import nltk
nltk.download('stopwords')
nltk.download('punkt')
nltk.download('wordnet')
# Download the missing punkt_tab resource
nltk.download('punkt_tab')

from nltk.stem import WordNetLemmatizer
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
import re

from spellchecker import SpellChecker
from contractions import fix

lemmatizer = WordNetLemmatizer()
stop_words = set(stopwords.words('english'))

def preprocess_text(text):
    # Lowercase the text
    text = text.lower()
    # Expand contractions
    text = fix(text)

    text = re.sub(r'<.*?>', '', text)
    text = re.sub(r'https?://\S+|www\.\S+', '', text)
    # Remove special characters (keeping words, spaces, hyphens, apostrophes)
    text = re.sub(r'^\W\s-',"", text)
    # Tokenize the text
    tokens = word_tokenize(text)
    # Remove stopwords and lemmatize
    tokens = [lemmatizer.lemmatize(word) for word in tokens if word not in stop_words]
    # Join tokens back into a single string
    return ' '.join(tokens)

X_train_tokens = [word_tokenize(text) for text in X_train]
X_test_tokens = [word_tokenize(text) for text in X_test]

```



```
# You also need to import Word2Vec
from gensim.models import Word2Vec

word2vec_model = Word2Vec(sentences=X_train_tokens, vector_size=100, window=5, min_count=1, workers=4)

# Function to average word vectors for a document
def document_vector(doc, model):
    # remove out-of-vocabulary words
    words = [word for word in doc if word in model.wv.index_to_key]
    if not words:
        return np.zeros(model.vector_size)
    return np.mean(model.wv[words], axis=0)

# Create document vectors for training and testing data
X_train_w2v = np.array([document_vector(doc, word2vec_model) for doc in X_train_tokens])
X_test_w2v = np.array([document_vector(doc, word2vec_model) for doc in X_test_tokens])

# Now X_train_w2v and X_test_w2v are defined and can be used for training the model

from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score

# Train Model
clf = LogisticRegression(max_iter=500)
# Use the newly created Word2Vec features for training
clf.fit(X_train_w2v, y_train)

# Predict on Test Set
# Use the newly created Word2Vec features for prediction
y_pred = clf.predict(X_test_w2v)

# Calculate Accuracy
accuracy = accuracy_score(y_test, y_pred)

print(f"Accuracy with Word2Vec features: {accuracy}")
```

```
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data] Package punkt is already up-to-date!
[nltk_data] Downloading package wordnet to /root/nltk_data...
[nltk_data] Package wordnet is already up-to-date!
[nltk_data] Downloading package punkt_tab to /root/nltk_data...
[nltk_data] Unzipping tokenizers/punkt_tab.zip.
-----
ValueError                                Traceback (most recent call last)
<ipython-input-58-4a55ae091498> in <cell line: 0>()
    41
    42 # You also need to import Word2Vec
--> 43 from gensim.models import Word2Vec
    44
    45
```

↕ 5 frames

/usr/local/lib/python3.11/dist-packages/gensim/_matutils.pyx in init gensim._matutils()

ValueError: numpy.dtype size changed, may indicate binary incompatibility. Expected 96 from C header, got 88 from PyObject

```
# Fit and Transform
X_train_counts = count_vectorizer.fit_transform(X_train)
X_test_counts = count_vectorizer.transform(X_test)
X_train_tfidf = tfidf_vectorizer.fit_transform(X_train)
X_test_tfidf = tfidf_vectorizer.transform(X_test)
```

Implementing LSTM (Long Short-Term Memory) for Sentiment Analysis on Textual Data

```
X_train, X_test, y_train, y_test = train_test_split(df['text'], df['label'])
```

```
y_train
```



	label
38125	NaN
2247	NaN
26908	NaN
29934	NaN
34861	NaN
...	...
38804	NaN
9727	NaN
8518	NaN
32495	NaN
1093	NaN

30000 rows × 1 columns

dtype: float64

```
import tensorflow as tf
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Embedding, LSTM, Dense, Dropout, Bidirectional
```

```
# Ensure necessary libraries are installed and updated
!pip install tensorflow
!pip install opencv-python
!pip install numpy --upgrade --force-reinstall
```

```
# Tokenization Parameters
max_vocab_size = 20000 # Maximum number of words to keep
max_sequence_length = 300 # Maximum length of each review
```

```
# Tokenize Text
tokenizer = Tokenizer(num_words=max_vocab_size, oov_token="<OOV>")
tokenizer.fit_on_texts(X_train)
```



```
Requirement already satisfied: tensorflow in /usr/local/lib/python3.11/dist-packages (2.18.0)
Requirement already satisfied: absl-py>=1.0.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (1.4.0)
Requirement already satisfied: astunparse>=1.6.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (1.6.3)
Requirement already satisfied: flatbuffers>=24.3.25 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (25.2.10)
Requirement already satisfied: gast!=0.5.0,!0.5.1,!0.5.2,>=0.2.1 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (0.6.0)
Requirement already satisfied: google-pasta>=0.1.1 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (0.2.0)
Requirement already satisfied: libclang>=13.0.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (18.1.1)
Requirement already satisfied: opt-einsum>=2.3.2 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (3.4.0)
Requirement already satisfied: packaging in /usr/local/lib/python3.11/dist-packages (from tensorflow) (24.2)
Requirement already satisfied: protobuf!=4.21.0,!4.21.1,!4.21.2,!4.21.3,!4.21.4,!4.21.5,<6.0.0dev,>=3.20.3 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (75.2.0)
Requirement already satisfied: requests<3,>=2.21.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (2.32.3)
Requirement already satisfied: setuptools in /usr/local/lib/python3.11/dist-packages (from tensorflow) (75.2.0)
Requirement already satisfied: six>=1.12.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (1.17.0)
Requirement already satisfied: termcolor>=1.1.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (3.1.0)
Requirement already satisfied: typing-extensions>=3.6.6 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (4.13.2)
Requirement already satisfied: wrapt>=1.11.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (1.17.2)
Requirement already satisfied: grpcio<2.0,>=1.24.3 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (1.71.0)
Requirement already satisfied: tensorboard<2.19,>=2.18 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (2.18.0)
Requirement already satisfied: keras>=3.5.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (3.8.0)
Requirement already satisfied: numpy<2.1.0,>=1.26.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (1.26.4)
Requirement already satisfied: h5py>=3.11.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (3.13.0)
Requirement already satisfied: ml-dtypes<0.5.0,>=0.4.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (0.4.1)
Requirement already satisfied: tensorflow-io-gcs-filesystem>=0.23.1 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (0.37.1)
Requirement already satisfied: wheel<1.0,>=0.23.0 in /usr/local/lib/python3.11/dist-packages (from astunparse>=1.6.0->tensorflow) (0.45.1)
Requirement already satisfied: rich in /usr/local/lib/python3.11/dist-packages (from keras>=3.5.0->tensorflow) (13.9.4)
Requirement already satisfied: namex in /usr/local/lib/python3.11/dist-packages (from keras>=3.5.0->tensorflow) (0.0.9)
Requirement already satisfied: optree in /usr/local/lib/python3.11/dist-packages (from keras>=3.5.0->tensorflow) (0.15.0)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.11/dist-packages (from requests<3,>=2.21.0->tensorflow) (3.10)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.11/dist-packages (from requests<3,>=2.21.0->tensorflow) (3.10)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.11/dist-packages (from requests<3,>=2.21.0->tensorflow) (2.4)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.11/dist-packages (from requests<3,>=2.21.0->tensorflow) (202)
Requirement already satisfied: markdown>=2.6.8 in /usr/local/lib/python3.11/dist-packages (from tensorboard<2.19,>=2.18->tensorflow) (3)
Requirement already satisfied: tensorboard-data-server<0.8.0,>=0.7.0 in /usr/local/lib/python3.11/dist-packages (from tensorboard<2.19,>=2.18->tensorflow) (0.8.0)
Requirement already satisfied: werkzeug>=1.0.1 in /usr/local/lib/python3.11/dist-packages (from tensorboard<2.19,>=2.18->tensorflow) (3)
Requirement already satisfied: MarkupSafe>=2.1.1 in /usr/local/lib/python3.11/dist-packages (from werkzeug>=1.0.1->tensorboard<2.19,>=2.18->tensorflow) (3)
Requirement already satisfied: markdown-it-py>=2.2.0 in /usr/local/lib/python3.11/dist-packages (from rich->keras>=3.5.0->tensorflow) (3)
```

```
Requirement already satisfied: pygments<3.0.0,>=2.13.0 in /usr/local/lib/python3.11/dist-packages (from rich->keras>=3.5.0->tensorflow)
Requirement already satisfied: mdurl~=0.1 in /usr/local/lib/python3.11/dist-packages (from markdown-it-py>=2.2.0->rich->keras>=3.5.0->te
Requirement already satisfied: opencv-python in /usr/local/lib/python3.11/dist-packages (4.11.0.86)
Requirement already satisfied: numpy>=1.21.2 in /usr/local/lib/python3.11/dist-packages (from opencv-python) (1.26.4)
Collecting numpy
  Using cached numpy-2.2.5-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (62 kB)
Using cached numpy-2.2.5-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (16.4 MB)
Installing collected packages: numpy
  Attempting uninstall: numpy
    Found existing installation: numpy 1.26.4
    Uninstalling numpy-1.26.4:
      Successfully uninstalled numpy-1.26.4
ERROR: pip's dependency resolver does not currently take into account all the packages that are installed. This behaviour is the source
gensim 4.3.3 requires numpy<2.0,>=1.18.5, but you have numpy 2.2.5 which is incompatible.
tsfresh 0.21.0 requires scipy>=1.14.0; python_version >= "3.10", but you have scipy 1.13.1 which is incompatible.
numba 0.60.0 requires numpy<2.1,>=1.22, but you have numpy 2.2.5 which is incompatible.
tensorflow 2.18.0 requires numpy<2.1.0,>=1.26.0, but you have numpy 2.2.5 which is incompatible.
Successfully installed numpy-2.2.5
```

```
X_train_seq = tokenizer.texts_to_sequences(X_train)
X_train_seq = tokenizer.texts_to_sequences(X_train)
X_test_seq = tokenizer.texts_to_sequences(X_test)
# Pad Sequences (Ensure all inputs are of the same length)
X_train_padded = pad_sequences(X_train_seq, maxlen=max_sequence_length, padding='post')
X_test_padded = pad_sequences(X_test_seq, maxlen=max_sequence_length, padding='post')
```

```
X_test_padded
```

```
array([[ 533, 1469, 23, ..., 0, 0, 0],
       [ 2, 877, 484, ..., 0, 0, 0],
       [ 11, 238, 494, ..., 0, 0, 0],
       ...,
       [ 12, 16, 7, ..., 0, 0, 0],
       [1753, 198, 30, ..., 0, 0, 0],
       [ 70, 11, 457, ..., 0, 0, 0]], dtype=int32)
```

Define LSTM model

```
from tensorflow.keras.layers import Embedding, LSTM, Dense, Dropout, Bidirectional, GRU
```

```
model = Sequential()
model.add(Embedding(input_dim=max_vocab_size, output_dim=128, input_length=max_sequence_length))
model.add(LSTM(128, dropout=0.2, recurrent_dropout = 0.2, return_sequences=True))
model.add(Bidirectional(GRU(64, return_sequences=False)))
model.add(Dense(1, activation = "sigmoid"))
```